

Differentiable Influence Minimization via Continuous Relaxation

Final Project Report: CMSC838B

Mihil Sreenilayam

University of Maryland, College Park

December 19, 2025

Motivation: The Cost of Connection

The Problem:

- Networks amplify diffusion: viruses, misinformation, and malware.
- **Influence Minimization (IMIN):**
Given a budget k , which k edges should we cut to maximally stifle the spread?

Motivation: The Cost of Connection

The Problem:

- Networks amplify diffusion: viruses, misinformation, and malware.
- **Influence Minimization (IMIN):**
Given a budget k , which k edges should we cut to maximally stifle the spread?

The Challenge:

- It is a *combinatorial* problem (discrete choices).
- It is a *stochastic* problem (random spread).
- **Result:** Standard gradients don't work.
We usually rely on slow heuristics or greedy algorithms.




Figure: Network diffusion often relies on critical "bridge" edges.

Theoretical Formulation

We model diffusion via the **Independent Cascade (IC)** model.

The Objective

Minimize expected spread $\sigma(G)$ subject to a budget constraint:

$$\min_{E' \subset E} \sigma(G(V, E')) \quad \text{s.t.} \quad |E| - |E'| \leq k$$

Theoretical Formulation

We model diffusion via the **Independent Cascade (IC)** model.

The Objective

Minimize expected spread $\sigma(G)$ subject to a budget constraint:

$$\min_{E' \subset E} \sigma(G(V, E')) \quad \text{s.t.} \quad |E| - |E'| \leq k$$

Why is this hard?

- ➊ **#P-Hardness:** Calculating $\sigma(G)$ exactly requires summing over $2^{|E|}$ possible live-edge realizations.
- ➋ **Non-Differentiable:** The spread function is discrete.

$\frac{\partial \sigma}{\partial \text{edge}_{uv}}$ is undefined.

The Solution: Differentiable Relaxation

We propose a two-stage differentiable pipeline to bypass these hurdles.

1. Surrogate Modeling

- Replace the expensive Monte Carlo simulator with a **Graph Neural Network (GNN)**.
- The GNN approximates the influence function differentiably.

The Solution: Differentiable Relaxation

We propose a two-stage differentiable pipeline to bypass these hurdles.

1. Surrogate Modeling

- Replace the expensive Monte Carlo simulator with a **Graph Neural Network (GNN)**.
- The GNN approximates the influence function differentiably.

2. Continuous Relaxation

- Replace binary edges $\{0, 1\}$ with continuous weights $w_{uv} \in [0, 1]$.
- Turn the discrete "switch" into a smooth "dimmer knob."

Goal: Allow gradients to flow from the predicted spread back to the graph structure.

Project Contributions

Relationship to Prior Work (DiffIM)

This project builds on the theoretical framework of Lee et al. (2025).

My Specific Contributions:

- ① **Reproducible Implementation:** Built the full pipeline in PyTorch Geometric (Surrogate + Relaxation Loop).
- ② **Feature Engineering:** Designed the structural feature set ($d = 3$) to enable size-invariant generalization.
- ③ **The "Lambda Sweep":** Developed an evaluation protocol to map the full Pareto frontier of sparsity vs. influence.
- ④ **Baseline Comparison:** Conducted rigorous testing against random baselines on synthetic (ER, BA, WS) and real (SNAP) graphs.

Methodology: GNN Surrogate Architecture

We utilize a Graph Convolutional Network (GCN) to map graph topology to a scalar influence score \hat{y} .

Input Features: Strictly structural ($x_v = [\text{deg}, \text{clust}, \text{neigh_deg}]$) to ensure inductive generalization.

Optimization: The Sparsity Penalty

To remove edges, we minimize a compound loss function:

$$\mathcal{L}(\mathcal{W}) = \underbrace{f_{\theta}(\mathcal{W})}_{\text{Influence Min.}} + \lambda \cdot \underbrace{\sum_{(u,v) \in E} |1 - w_{uv}|}_{\text{Sparsity Penalty}}$$

- **First Term:** Pushes weights to minimize spread (flatten the curve).
- **Second Term:** Penalizes deviation from the original graph (cost of removal).
- **λ (Lambda):** The "aggressiveness" knob.
 - High $\lambda \rightarrow$ Keep edges (Conservative).
 - Low $\lambda \rightarrow$ Remove edges (Aggressive).

Experimental Setup

- **Dataset:** $\approx 1,400$ graphs.
 - *Training:* Mix of Erdős–Rényi, Barabási–Albert, Watts–Strogatz.
 - *Testing:* Held-out synthetic graphs + Real-world SNAP subgraphs (Email-Eu-core).
- **Ground Truth:** Monte Carlo Simulations (500 runs per graph).
- **Baseline: Random Edge Removal.**
 - Crucial: We match the *exact number of edges removed* to ensure fair comparison.

Results: Surrogate Fidelity

Can a neural network actually predict stochastic influence? Yes.

Graph Type	Pearson r	Spearman ρ
Erdős–Rényi	0.982	0.975
Watts–Strogatz	0.971	0.966
Barabási–Albert	0.954	0.941
Overall	0.969	0.960

Table: Correlation between GNN prediction and Monte Carlo truth on test set.

The high rank correlation (ρ) confirms the surrogate correctly identifies which graph states are "worse" than others.

Results: Sparsity-Influence Trade-off

Analysis:

- **Orange (Random):** Linear/Convex decay. Requires massive deletion to stop spread.
- **Blue (Ours):** Concave decay. Significant influence drop with minimal deletion.
- **Interpretation:** The gradient method identifies *structural bottlenecks* (bridges) that random deletion misses.




Figure: Influence vs. Edge Retention Rate.

Limitations and Future Work

Current Limitations:

- **Generalization Gap:** Performance drops on real-world graphs ($r \approx 0.65$) due to complex community modularity not present in synthetic training data.
- **Scalability:** Full-batch GCN is memory intensive for massive graphs.

Future Directions:

- **Domain Adaptation:** Few-shot fine-tuning of the surrogate on real-world subgraphs.

Summary

- ① We formulated Influence Minimization as a **differentiable learning problem**.
- ② We implemented a pipeline using a **GCN Surrogate** and **Continuous Edge Relaxation**.
- ③ We demonstrated that gradient-based pruning significantly outperforms random baselines, identifying critical structural weaknesses in networks.
- ④ The framework provides a bridge between combinatorial graph theory and modern deep learning optimization.

Thank You! Questions?

References I

-  Lee, J., et al. (2025). *DiffIM: Differentiable Influence Minimization with Surrogate Modeling*. arXiv:2502.01031.
-  Kipf, T. N., & Welling, M. (2017). *Semi-Supervised Classification with Graph Convolutional Networks*. ICLR.
-  Kempe, D., Kleinberg, J., & Tardos, É. (2003). *Maximizing the spread of influence through a social network*. KDD.
-  Manchanda, S., et al. (2020). *Gcomb: Learning budget-constrained combinatorial algorithms over billion-sized graphs*. NeurIPS.